

MPI Non-blocking Communication



ICHEC
Irish Centre for High-End Computing



An Roinn Post, Fiontar agus Nuálaíochta
Department of Jobs, Enterprise and Innovation



AN ROINN
OIDEACHAIS AGUS SCILEANNA
DEPARTMENT OF
EDUCATION AND SKILLS

HEA
Higher Education Authority
an tÚdarás um Ard-Oideachas


www.ichec.ie

Deadlock

- Code in each MPI process:

```
MPI_Ssend(..., right_rank, ...)
```

```
MPI_Recv( ..., left_rank, ...)
```



- Same problem with standard send mode (MPI_Send), if MPI implementation chooses synchronous protocol

Non-Blocking Communications

- Separate communication into three phases:
- Initiate non-blocking communication
 - returns Immediately
 - routine name starting with MPI_I...
- Do some work
 - “latency hiding”
- Wait for non-blocking communication to complete

Non-Blocking Examples

- Non-blocking **send**

MPI_Isend(...)

doing some other work

MPI_Wait(...)




- Non-blocking **receive**

MPI_Irecv(...)

doing some other work


MPI_Wait(...)



 = waiting until operation locally completed


Non-Blocking Send

- Initiate non-blocking send
 - in the ring example: Initiate non-blocking send to the right neighbor
- Do some work:
 - in the ring example: Receiving the message from left neighbor
- Now, the message transfer can be completed
- Wait for non-blocking send to complete ✓



Non-Blocking Receive

- Initiate non-blocking receive
 - in the ring example: Initiate non-blocking receive from left neighbor
- Do some work:
 - in the ring example: Sending the message to the right neighbor
- Now, the message transfer can be completed
- Wait for non-blocking receive to complete



Non-blocking Synchronous Send

- C:
 - `MPI_Issend(buf, count, datatypeHandle, dest, tag, comm, OUT &requestHandle);`
 - `MPI_Wait(INOUT &requestHandle, &status);`
- Fortran:
 - call `MPI_Issend(buf, count, datatypeHnadle, dest, tag, comm, OUT requestHandle, ierror)`
 - call `MPI_Wait(INOUT requestHandle, status, ierror)`
- Request handle must be stored in local variables
 - C: `MPI_Request`
 - Fortran: integer
- buf must not be used between Issend and Wait (in all progr. languages)
- “Issend + Wait directly after Issend” is equivalent to blocking call (Ssend)

Non-blocking Receive

- C:
 - `MPI_Irecv(buf, count, datatypeHandle, source, tag, comm, OUT &requestHandle);`
 - `MPI_Wait(INOUT &requestHandle, &status);`
- Fortran:
 - call `MPI_Irecv (buf, count, datatype, source, tag, comm, OUT requestHandle, ierror)`
 - call `MPI_Wait(INOUT requestHandle, status, ierror)`
- buf must not be used between Irecv and Wait (in all progr. languages)

Blocking and Non-Blocking

- Send and receive can be blocking or non-blocking.

| Send Mode | Blocking Function | Nonblocking Function |
|-------------|-------------------|----------------------|
| Standard | MPI_Send | MPI_Isend |
| Synchronous | MPI_Ssend | MPI_Issend |
| Ready | MPI_Rsend | MPI_Irsend |
| Buffered | MPI_Bsend | MPI_Ibsend |

Completion

- C:
 - `MPI_Wait(&requestHandle, &status);`
 - `MPI_Test(&requestHandle, &flag, &status);`
- Fortran:
 - call `MPI_Wait(requestHandle, status, ierror)`
 - call `MPI_Test(requestHandle, flag, status, ierror)`
- one must
 - wait or
 - loop with TEST until request is completed, i.e., `flag == 1` or `.true.`

Multiple Non-Blocking Communications

- You have several request handles:
- Wait or test for completion of one message
 - `MPI_Waitany` / `MPI_Testany`
- Wait or test for completion of all messages
 - `MPI_Waitall` / `MPI_Testall`
- Wait or test for completion of as many messages as possible
 - `MPI_Waitsome` / `MPI_Testsome`

