

# Defining Communicators



**ICHEC**  
Irish Centre for High-End Computing



An Roinn Post, Fiontar agus Nuálaíochta  
Department of Jobs, Enterprise and Innovation



AN ROINN  
OIDEACHAIS AGUS SCILEANNA  
DEPARTMENT OF  
EDUCATION AND SKILLS

**HEA**  
Higher Education Authority  
in tUdaráis um Ard-Oideachas

[www.ichec.ie](http://www.ichec.ie)


# Communicators

- A communicator determines the “communication universe”
  - Source and destination of a message is identified by the process rank within the communicator
  - MPI\_COMM\_WORLD used so far
- Processes can be divided into subcommunicators also known as groups
  - Task level parallelism with process ‘groups’ performing separate duties together
  - Parallel I/O
  - Scalability → avoids unnecessary synchronization

# Motivations

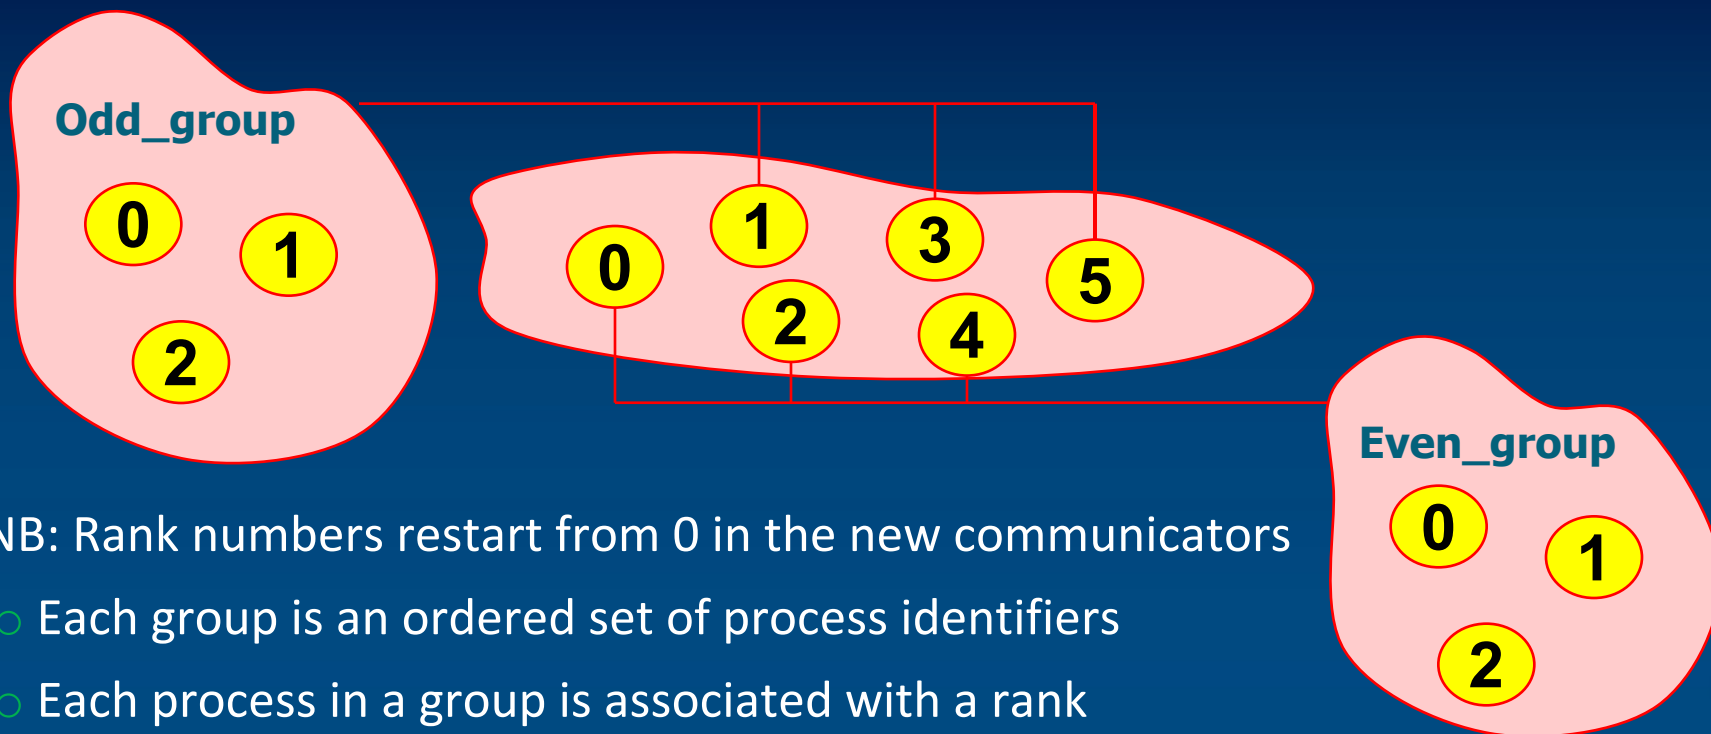
- Need to create sets of processes
  - For programming convenience
  - Make use of collective routines
- Need to map the abstract topology onto the natural topology of the problem domain
  - For programming convenience
  - For performance

# Working with groups



- Select processes ranks to create groups
- Associate to these groups *new* communicators
- Use these new communicators as usual
- `MPI_Comm_group(comm, group)` returns in *group* the group associated to the communicator *comm*
- Communicators are dynamic and can be created and destroyed during runtime
- A task can belong simultaneously to several communicators, and in each, has a unique ID

# Working with groups



- NB: Rank numbers restart from 0 in the new communicators
  - Each group is an ordered set of process identifiers
  - Each process in a group is associated with a rank
- WORLD {1, 3, 5} → Odd\_group {0, 1, 2}
- WORLD {0, 2, 4} → Even\_group {0, 1, 2}

# Method 1: MPI\_comm\_group

```
Odd_ranks={1, 3, 5}, Even_ranks={0, 2, 4};  
MPI_comm_group(MPI_COMM_WORLD, Group_all);  
MPI_Group_incl(Group_all, 3, Odd_ranks, &Odd_group);  
MPI_Group_incl(Group_all, 3, Even_ranks, &Even_group);  
int MPI_Comm_create(MPI_COMM_WORLD, Odd_group, Odd_Comm);  
int MPI_Comm_create(MPI_COMM_WORLD, Even_group, Even_Comm);
```

# Method 2: MPI\_Comm\_split

- Create new communicators based on 'colors' and 'keys'

`MPI_Comm_split(comm, color, key, &newcomm)`

<b>comm</b>	Communicator handle
<b>color</b>	Control of subset assignment, processes with the same color belong to the same new communicator
<b>key</b>	Control of rank assignment
<b>newcomm</b>	New communicator handle

- If `color = MPI_UNDEFINED`, a process does not belong to any of the new communicators

# Creating a communicator

```
if (myid%2 == 0) {
    color = 1;
} else {
    color = 2;
}

MPI_Comm_split(MPI_COMM_WORLD, color, myid, &subcomm);

MPI_Comm_rank(subcomm, %mysubid);

printf("I am rank %d in MPI_COMM_WORLD, but %d in Comm %d.\n", myid,
mysubid, color);
```

```
I am rank 2 in MPI_COMM_WORLD, but 1 in Comm 1.
I am rank 1 in MPI_COMM_WORLD, but 1 in Comm 2.
I am rank 4 in MPI_COMM_WORLD, but 1 in Comm 1.
I am rank 5 in MPI_COMM_WORLD, but 1 in Comm 2.
I am rank 3 in MPI_COMM_WORLD, but 1 in Comm 2.
I am rank 0 in MPI_COMM_WORLD, but 1 in Comm 1.
```





# Group & Communicator Management

Group	Communicator
<b>Accessors</b>	
<code>MPI_Group_size(...)</code>	<code>MPI_Comm_size(...)</code>
<code>MPI_Group_rank(...)</code>	<code>MPI_Comm_rank(...)</code>
...	...
<b>Constructors</b>	
<code>MPI_Comm_group(...)</code>	<code>MPI_Comm_create(...)</code>
<code>MPI_Group_incl(...)</code>	<code>MPI_Comm_split(...)</code>
<code>MPI_GROUP_EXCL(...)</code>	
...	
<b>Destructors</b>	
<code>MPI_Group_free(group)</code>	<code>MPI_Comm_free(comm)</code>