

# User Defined MPI Datatypes



# ICHEC

Irish Centre for High-End Computing



An Roinn Post, Fiontar agus Nuálaíochta  
Department of Jobs, Enterprise and Innovation



AN ROINN  
OIDEACHAIS AGUS SCILEANNA  
DEPARTMENT OF  
EDUCATION AND SKILLS

HEA  
Higher Education Authority  
an tÚdarás um Ard-Oideachas

  
[www.ichec.ie](http://www.ichec.ie)

# Data Layout and the Describing Datatype Handle

```
struct buff_layout
{
  int
  i_val[3];
  double d_val[5];
} buffer;
```

**Compiler**

```
array_of_types[0] = MPI_INT;
array_of_blocklengths[0] = 3;
array_of_displacements[0] = 0;
array_of_types[1] = MPI_DOUBLE;
array_of_blocklengths[1] = 5;
array_of_displacements[1] = ...;

MPI_Type_struct(2,
  array_of_blocklengths,
  array_of_displacements,
  array_of_types, &buff_datatype);

MPI_Type_commit(&buff_datatype);
```

`MPI_Send(&buffer, 1, buff_datatype, ...)`

**&buffer = the start  
address of the data**

**the datatype handle  
describes the data layout**



# Derived Datatypes — Type Maps

- A derived datatype is logically a pointer to a list of entries:
  - *basic datatype at displacement*

basic datatype 0	displacement of datatype 0
basic datatype 1	displacement of datatype 1
...	...
basic datatype n-1	displacement of datatype n-1

# Derived Datatypes — Type Maps

Example:




derived datatype handle

basic datatype	displacement
MPI_CHAR	0
MPI_INT	4
MPI_INT	8
MPI_DOUBLE	16

A derived datatype describes the memory layout of, e.g.,  
 structures,  
 common blocks,  
 subarrays,  
 some variables in the memory

# Contiguous Data

- The simplest derived datatype
- Consists of a number of contiguous items of the same datatype




- C: 

```
int MPI_Type_contiguous(int count, MPI_Datatype oldtype,
                        MPI_Datatype *newtype)
```
- Fortran: 

```
MPI_Type_contiguous(count, oldtype, newtype, ierror)
integer :: count, oldtype
integer :: newtype, ierror
```

# Vector Datatype



- C: `int MPI_Type_vector(int count, int blocklength, int stride, MPI_Datatype oldtype, MPI_Datatype *newtype)`
- Fortran: `MPI_Type_vector(count, blocklength, stride, oldtype, newtype, ierror)`  
`integer :: count, blocklength, stride`  
`integer :: oldtype, newtype, ierror`

# Sending a row using MPI\_TYPE\_vector

- C

- `MPI_Type_vector(1, 5, 1, MPI_INT, ARR_ROW)`


- Fortran

- `MPI_Type_vector(5, 1, 4, MPI_INT, ARR_ROW)`

`MPI_Type_Commit(ARR_ROW)`


`MPI_Send(&buf ..., ARR_ROW...)`

`MPI_Recv(&buf ..., ARR_ROW...)`



# Sending a column using MPI\_TYPE\_vector


- C
  - `MPI_Type_vector(4, 1, 5, MPI_INT, ARR_COL)`
- Fortran
  - `MPI_Type_vector(1, 4, 1, MPI_INT, ARR_COL)`
- `MPI_Type_Commit(ARR_COL)`
- `MPI_Send(buf ..., ARR_COL...)`
- `MPI_Recv(buf ..., ARR_COL...)`



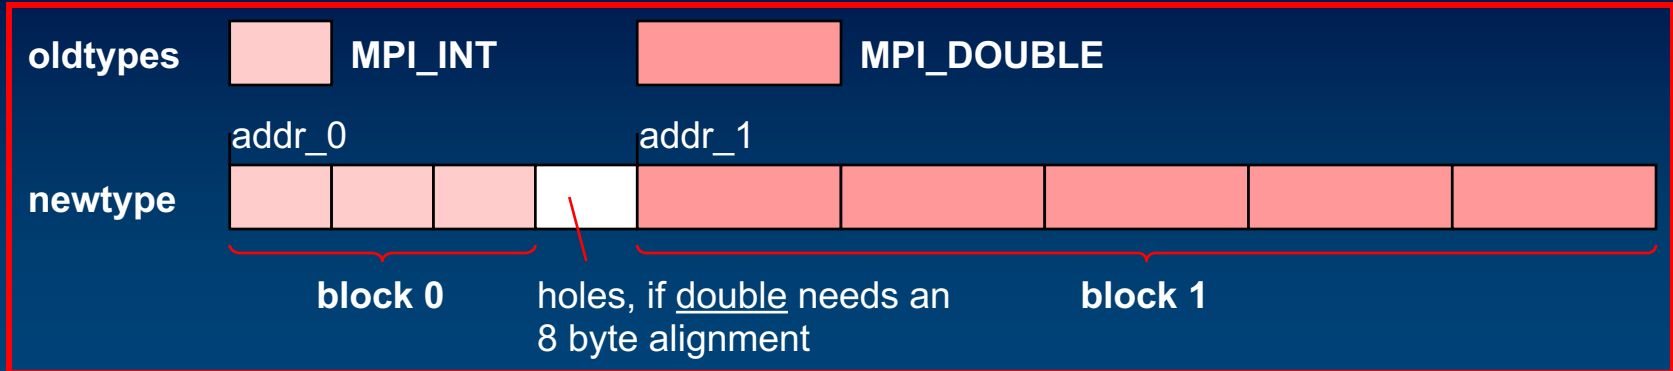


# Sending a sub-matrix using MPI\_TYPE\_vector

- C
  - `MPI_Type_vector(2, 3, 5, MPI_INT, SUBMAT)`
- Fortran
  - `MPI_Type_vector(3, 2, 4, MPI_INT, SUBMAT)`
- `MPI_Type_Commit(SUBMAT)`
- `MPI_Send(&buf ..., SUBMAT...)`
- `MPI_Recv(&buf ..., SUBMAT...)`



# Struct Datatype



- C: `int MPI_Type_struct(int count, int *array_of_blocklengths, MPI_Aint *array_of_displacements, MPI_Datatype *array_of_types, MPI_Datatype *newtype)`
- Fortran: `MPI_Type_struct(count, array_of_blocklengths, array_of_displacements, array_of_types, newtype, ierror)`

```

count = 2
array_of_blocklengths = ( 3,      5      )
array_of_displacements = ( 0,      addr_1 - addr_0 )
array_of_types = ( MPI_INT, MPI_DOUBLE )
  
```

# How to compute the displacement

- $\text{array\_of\_displacements}[i] := \text{address}(\text{block}_i) - \text{address}(\text{block}_0)$
- MPI-1
  - C: `int MPI_Address(void* location, MPI_Aint *address)`
  - Fortran: `MPI_ADDRESS(location, address, ierror)`  

`<type>            location(*)`  
`integer ::       address, ierror`

# Committing a Datatype

- Before a datatype handle is used in message passing communication,  
**it needs to be committed with MPI\_TYPE\_COMMIT.**
- This must be done only once.
- C: `int MPI_Type_commit(MPI_Datatype *datatype);`
- Fortran: `MPI_TYPE_COMMIT(datatype, ierror)`  
`integer :: datatype, ierror`

IN-OUT argument


# Size and Extent of a Datatype

- Size := number of bytes that have to be transferred.
- Extent := spans from first to last byte.
- Basic datatypes: Size = Extent = number of bytes used by the compiler.
- Derived datatypes, an example:

oldtype



newtype



size := 6 \* size(oldtype)

extent := 8 \* extent(oldtype)

better visualization of newtype:



# Size and Extent of a Datatype

- MPI-1:

- C: `int MPI_Type_size(MPI_Datatype datatype, int *size)`

- `int MPI_Type_extent(MPI_Datatype datatype, MPI_Aint *extent)`

- Fortran: `MPI_Type_size(datatype, size, ierror)`

- `integer :: datatype, size, ierror`

- `MPI_Type_extent(datatype, extent, ierror)`

- `integer :: datatype, extent, ierror`